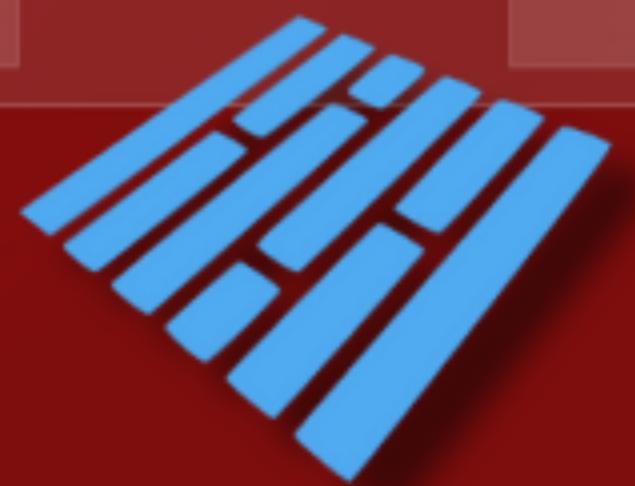




# INTRODUCTION TO Parquet



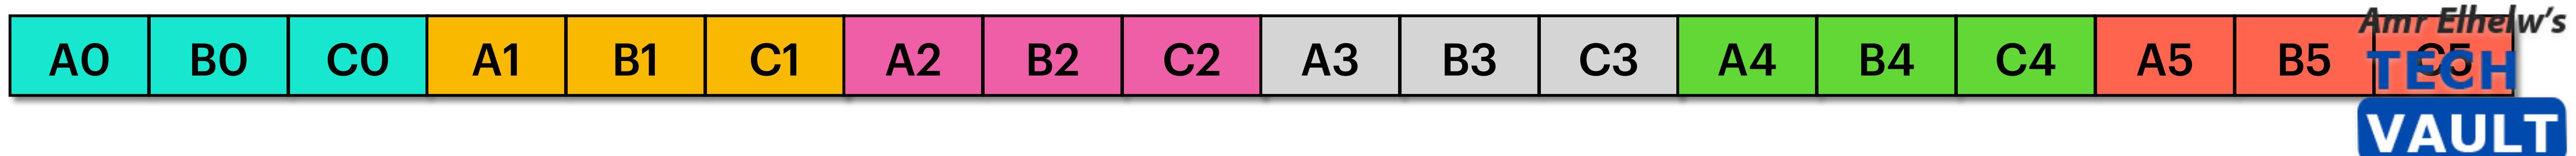
AMR ELHELW

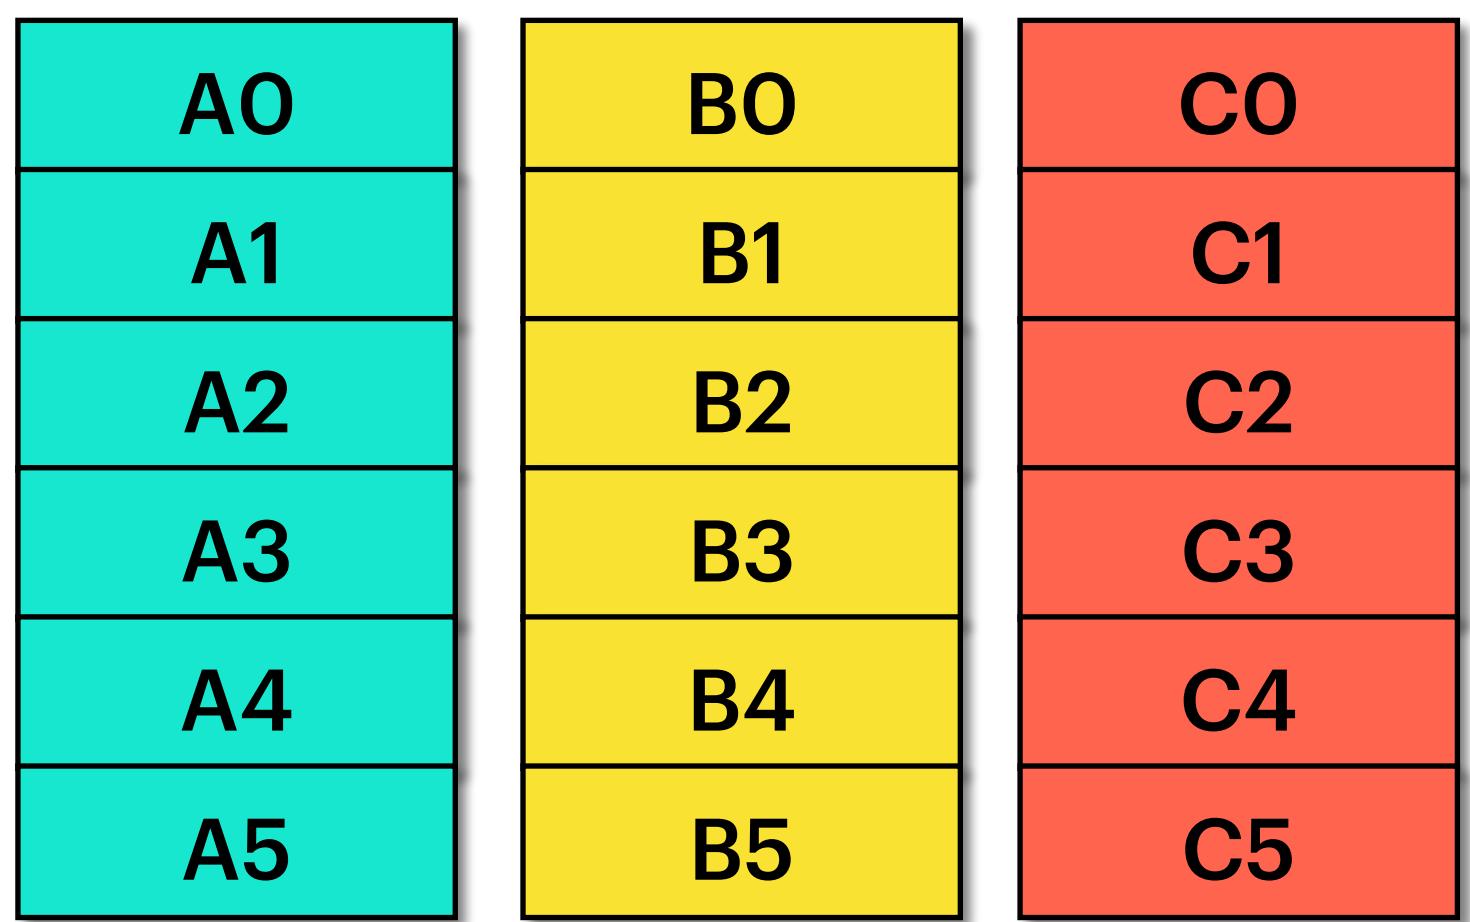
	Col A	Col B	Col C
Row 0	A0	B0	C0
Row 1	A1	B1	C1
Row 2	A2	B2	C2
Row 3	A3	B3	C3
Row 4	A4	B4	C4
Row 5	A5	B5	C5

A0	B0	C0
A1	B1	C1
A2	B2	C2
A3	B3	C3
A4	B4	C4
A5	B5	C5

Row-oriented

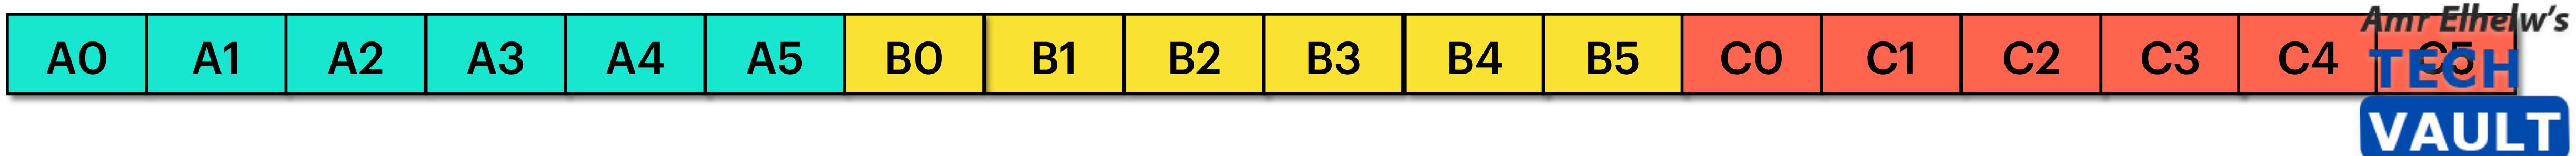
## Physical Storage





## Column-oriented (Columnar)

### Physical Storage



## OLTP

- Access a *small number of rows*
- *Read/Update several attributes from the same row.*
- *Primary focus is operational*

```
select *  
from employees  
where dept = 'IT'
```

## OLAP

- Read a *lot of rows (usually very large tables)*
- *Focus on a small number of columns in order to do some kind of aggregation*
- *Primary focus here is analytics*

```
select avg(salary)  
from employees
```

```
select *\nfrom employees\nwhere dept = 'IT'
```

A0	B0	C0
A1	B1	C1
A2	B2	C2
A3	B3	C3
A4	B4	C4
A5	B5	C5

## Row-oriented

OLTP

A0	B0	C0
A1	B1	C1
A2	B2	C2
A3	B3	C3
A4	B4	C4
A5	B5	C5

## Columnar

OLTP

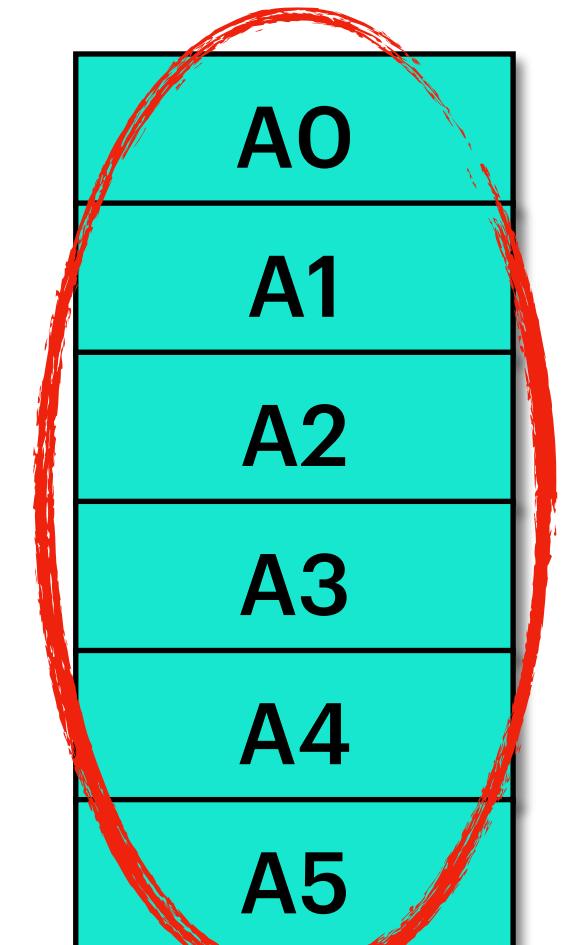
```
select avg(salary)  
from employees
```

A0	B0	C0
A1	B1	C1
A2	B2	C2
A3	B3	C3
A4	B4	C4
A5	B5	C5

**Row-oriented**

**OLTP** 

**OLAP** 



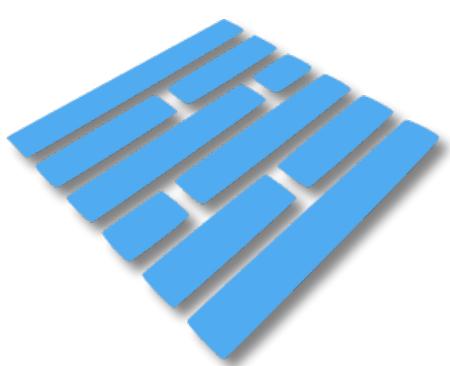
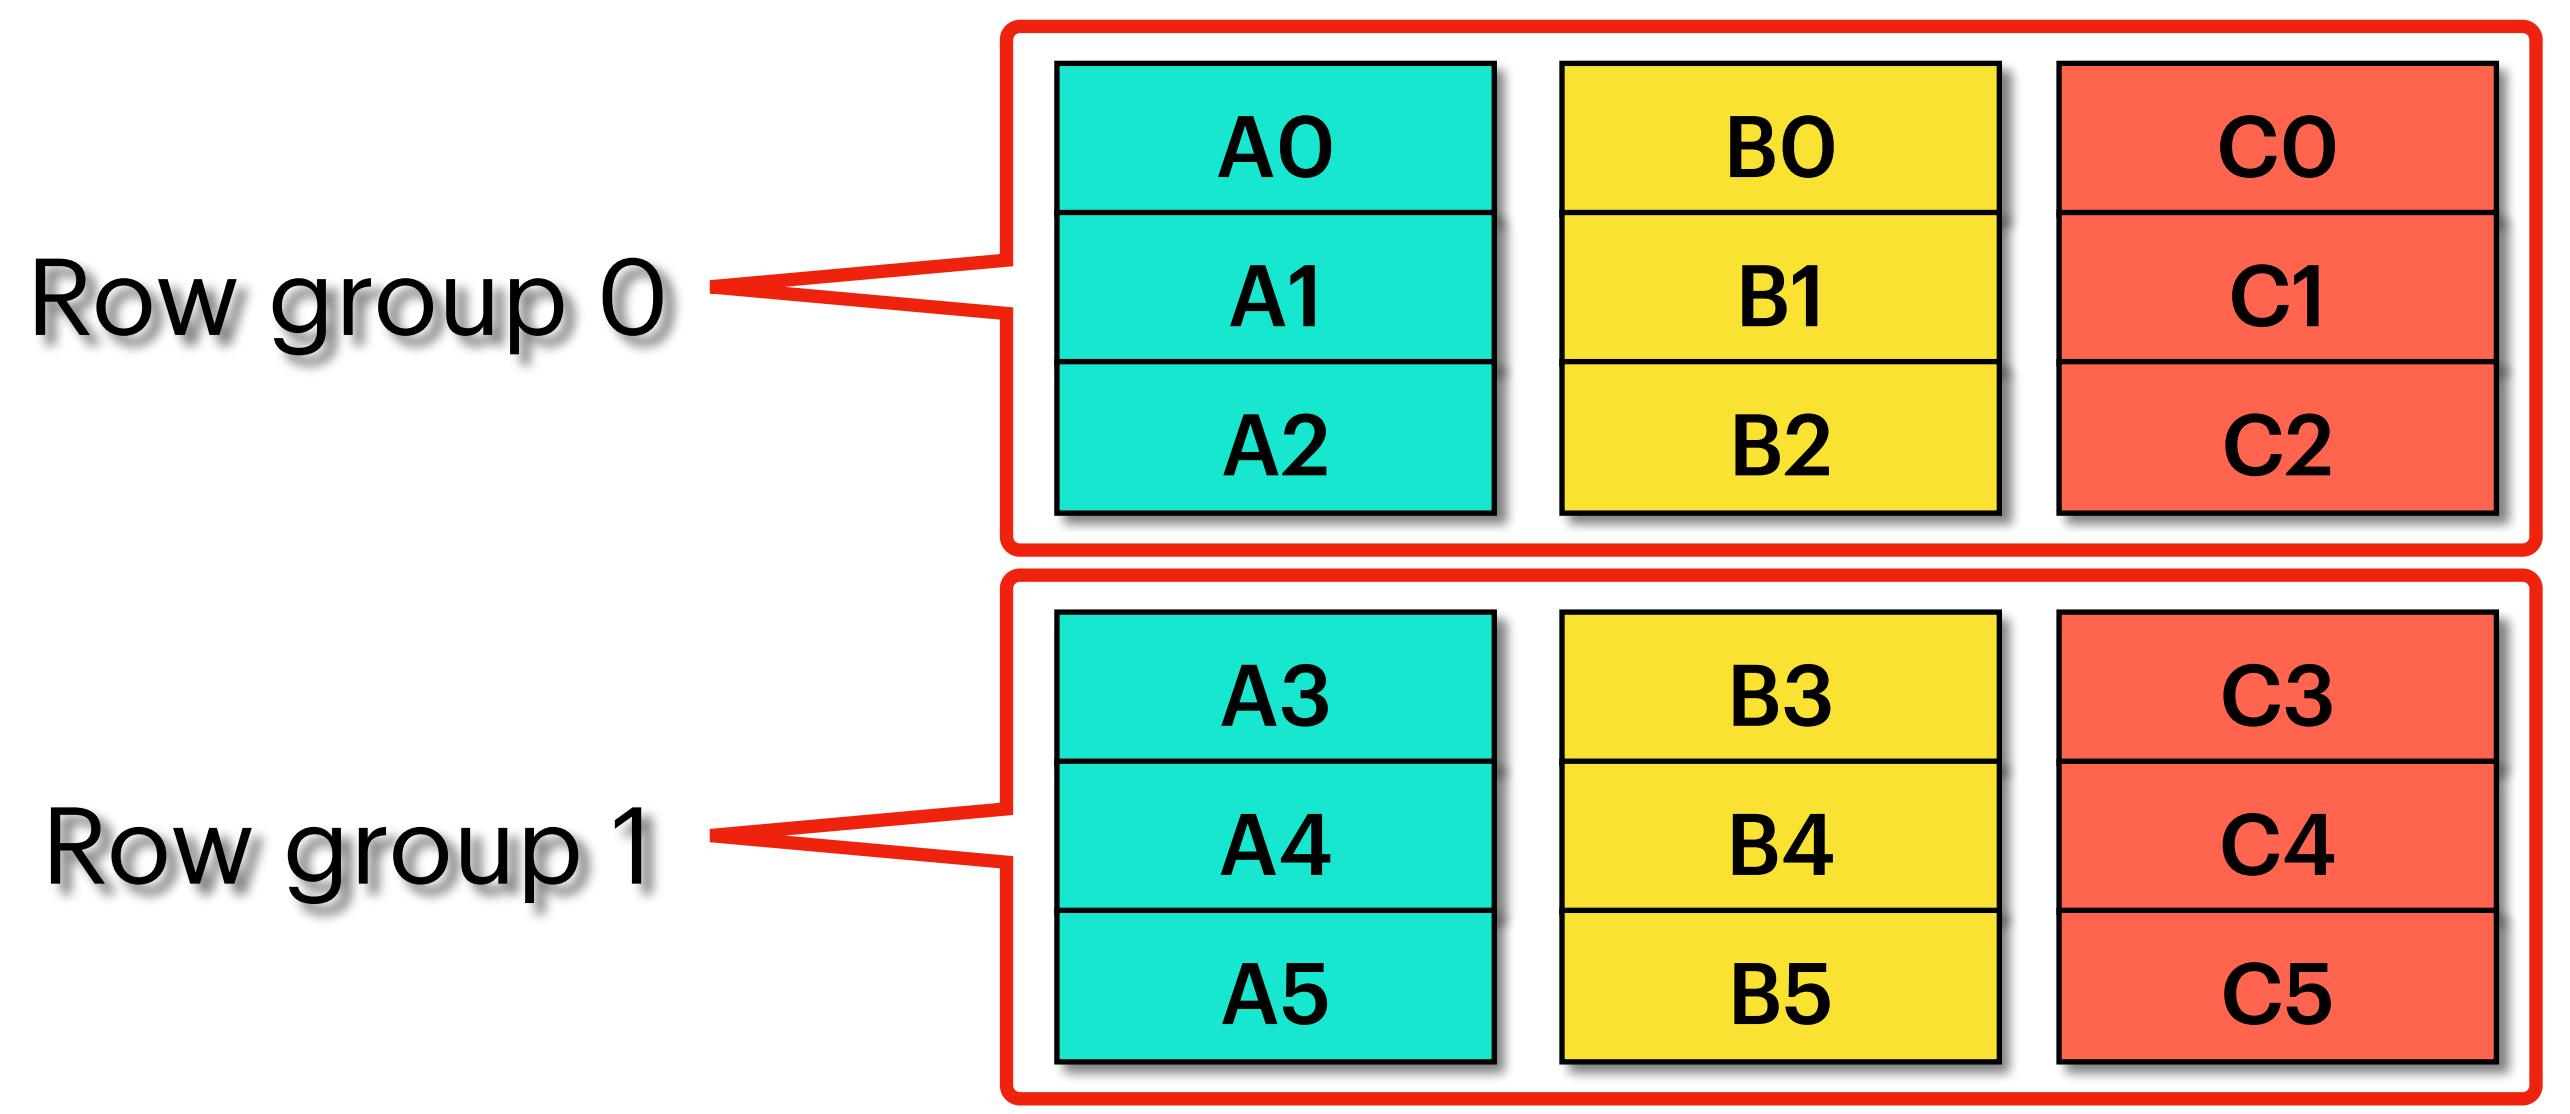
A0	BO	CO
A1	B1	C1
A2	B2	C2
A3	B3	C3
A4	B4	C4
A5	B5	C5

BO	CO
B1	C1
B2	C2
B3	C3
B4	C4
B5	C5

**Columnar**

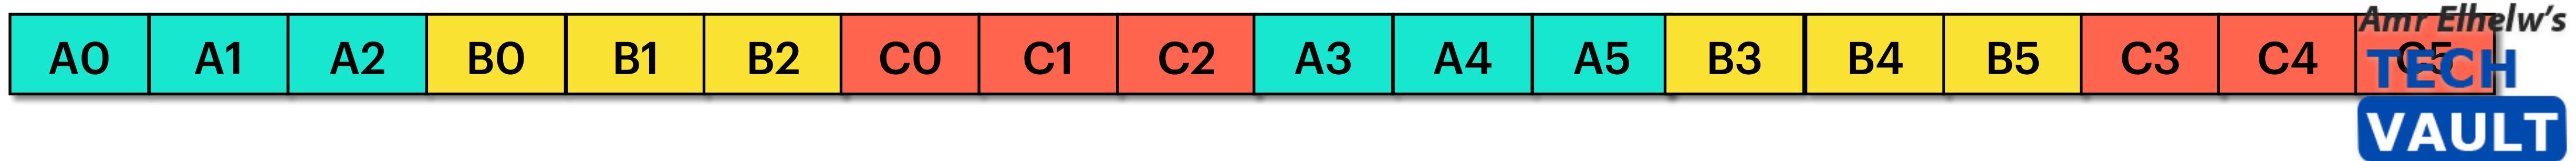
**OLTP** 

**OLAP** 



# Parquet

## Physical Storage





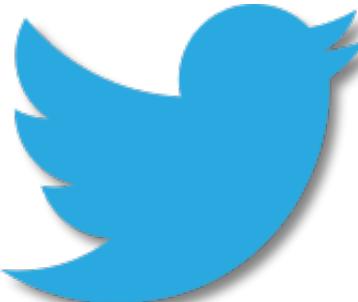
<https://parquet.apache.org/docs/overview>

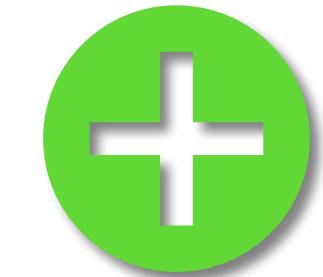
[Documentation](#) / Overview

# Overview

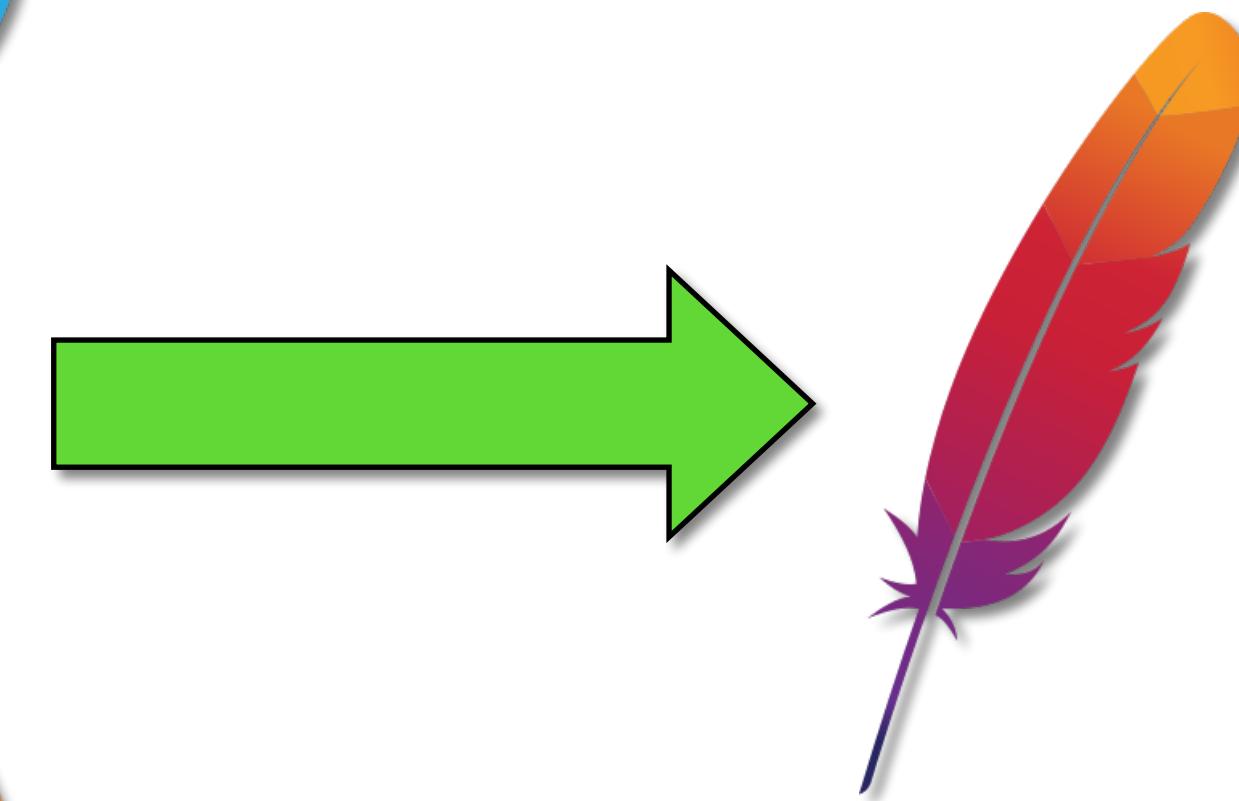
All about Parquet.

Apache Parquet is an **open source**, **column-oriented data file format** designed for **efficient data storage and retrieval**. It provides high performance **compression and encoding** schemes to handle **complex data** in bulk and is supported in **many programming language and analytics tools**.

twitter 



CLOUDERA



THE **APACHE**<sup>®</sup>  
SOFTWARE FOUNDATION

## Parquet File

### Row Group

Column  
Chunk

Column  
Chunk

Column  
Chunk

...

Column  
Chunk

### Row Group

Column  
Chunk

Column  
Chunk

Column  
Chunk

...

Column  
Chunk

⋮

### Row Group

Column  
Chunk

Column  
Chunk

Column  
Chunk

...

Column  
Chunk

Footer

Column Chunk

Page

Page

Page

Page

Page

Page

# Footer Contents

- Schema (Column names, types, etc.)
- Total number of rows in the file
- **Row group metadata (for each Row Group)**
  - # rows, byte size, file offset
  - **Column Chunk metadata (for each Column Chunk)**
    - File offset, chunk size
    - # distinct values in chunk, # nulls in chunk
    - Min/max values in chunk
    - Encoding/compression info for each page

## Parquet File - Table: Employee

Row Group 0

Column  
Chunk  
"id"

Column  
Chunk  
"Name"

Column  
Chunk  
"Dept"

Column  
Chunk  
"Salary"

Row Group 1

Column  
Chunk  
"id"

Column  
Chunk  
"Name"

Column  
Chunk  
"Dept"

Column  
Chunk  
"Salary"

Row Group 2

Column  
Chunk  
"id"

Column  
Chunk  
"Name"

Column  
Chunk  
"Dept"

Column  
Chunk  
"Salary"

Row Group 3

Column  
Chunk  
"id"

Column  
Chunk  
"Name"

Column  
Chunk  
"Dept"

Column  
Chunk  
"Salary"

Footer

```
select avg(salary)  
from employees
```

## Parquet File - Table: Employee

Row Group 0

Column  
Chunk  
"id"

Column  
Chunk  
"Name"

Column  
Chunk  
"Dept"

Column  
Chunk  
"Salary"

Row Group 1

Column  
Chunk  
"id"

Column  
Chunk  
"Name"

Column  
Chunk  
"Dept"

Column  
Chunk  
"Salary"

Row Group 2

Column  
Chunk  
"id"

Column  
Chunk  
"Name"

Column  
Chunk  
"Dept"

Column  
Chunk  
"Salary"

Row Group 3

Column  
Chunk  
"id"

Column  
Chunk  
"Name"

Column  
Chunk  
"Dept"

Column  
Chunk  
"Salary"

Footer

```
select avg(salary)  
from employee  
where id between 400 and 600
```

**Row group 0:**

"id": min = 102, max = 256

"dept": min = 10, max = 55

...

**Row group 1:**

"id": min = 260, max = 529

"dept": min = 5, max = 48

...

**Row group 2:**

"id": min = 570, max = 863

"dept": min = 18, max = 57

...

**Row group 3:**

"id": min = 869, max = 1104

"dept": min = 12, max = 41

...



## Parquet File - Table: Employee

Row Group 0

Column  
Chunk  
"id"

Column  
Chunk  
"Name"

Column  
Chunk  
"Dept"

Column  
Chunk  
"Salary"

Row Group 1

Column  
Chunk  
"id"

Column  
Chunk  
"Name"

Column  
Chunk  
"Dept"

Column  
Chunk  
"Salary"

Row Group 2

Column  
Chunk  
"id"

Column  
Chunk  
"Name"

Column  
Chunk  
"Dept"

Column  
Chunk  
"Salary"

Row Group 3

Column  
Chunk  
"id"

Column  
Chunk  
"Name"

Column  
Chunk  
"Dept"

Column  
Chunk  
"Salary"

Footer

```
select avg(salary)  
from employee  
where dept = 30
```

**Row group 0:**

"id": min = 102, max = 256

"dept": min = 10, max = 55

...

**Row group 1:**

"id": min = 260, max = 529

"dept": min = 5, max = 48

...

**Row group 2:**

"id": min = 570, max = 863

"dept": min = 18, max = 57



...

**Row group 3:**

"id": min = 869, max = 1104

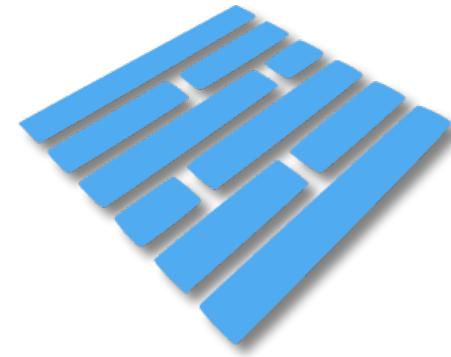
"dept": min = 12, max = 41



...

# Note on file size

- **Many small files**
  - Overhead of fetching multiple files
  - Setting up reader objects
  - Too many footers to read
- **Few huge files**
  - Footer gets too large — too many row groups, with info for each
- **No “magic number” for how big a file should be**
  - Trial and error
  - System specific guidelines

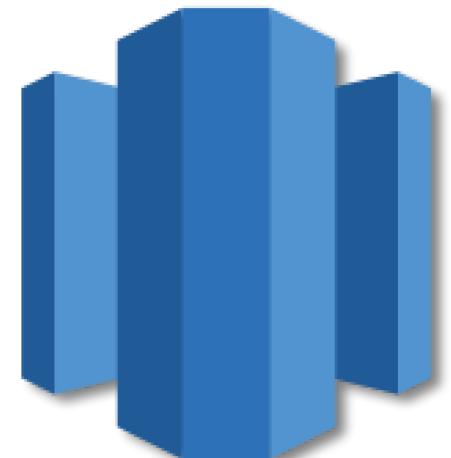


# Parquet

SQL



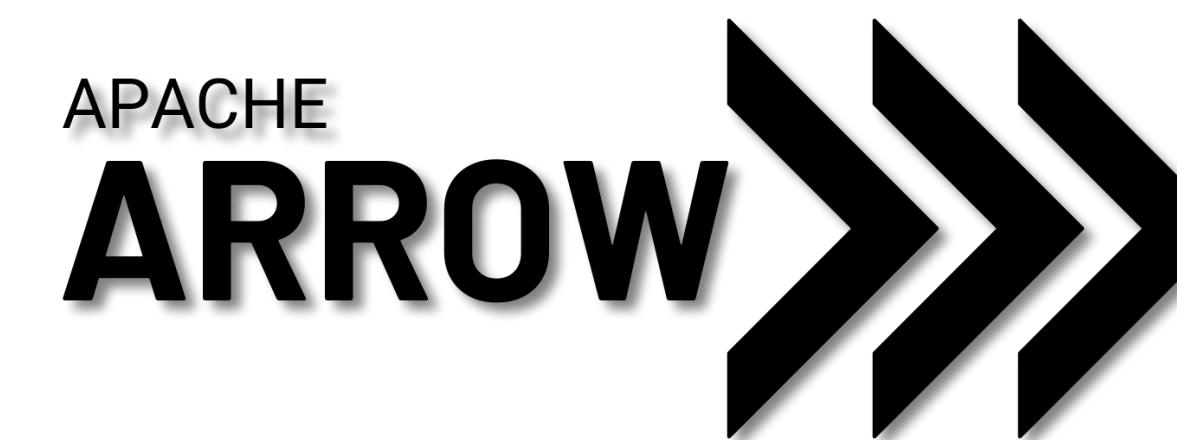
Google  
Big Query



amazon  
REDSHIFT



Programming Languages



C, C++, C#, Go,  
Java, JavaScript,  
Python, R, Ruby,  
Rust ...

Python

# Limitations

- Overhead for small datasets
  - Metadata adds cost
- Not ideal for frequent appends/updates
  - Requires rewrite
- Not suited for OLTP workloads
  - Inefficient for Row-Level Operations